

Software Project Day

Group 4: Matthew Witte, Jean Cheng, David Wilson

Assumptions

- time between questions for a team
- average lunch break
- arrival and departure times
- question after 4:00 pm manager doesn't answer them
- from 9:00 am to 2:59 pm there's a linear probability that any of the developer can get lunch
- manager has exactly one hour for lunch
- can't answer questions during lunch
- questions aren't answered in a first come/first served fashion
- stat for time working includes meetings, time waited for manager, but doesn't include lunch
- lunch lasts exactly one hour for the manager
- the manager doesn't necessarily arrive on time to their 10:00 am and 2:00 pm meetings
- at any time of interval there is a chance for people to ask questions

Document results of experimenting with different assumptions

Design Decisions

Data Structures

We chose to create a Clock class for representing the passage of time, so that it could be handled on a single thread and accessed by the employees and manager. Because LocalTime is only available in Java 8, we created a simple Time class nested within Clock to represent each moment. We broke encapsulation of the individual hour and minute values for ease of access, but the clock only ever provides a copy of its private Time object tracking the current time.

A `ConferenceRoom` class is used in a basic shared resource capacity, with a single boolean lock ensuring only one meeting is being held in the room at a time.

The `Employee` class handles all of the employees meetings and lunch. These meetings and lunch time are subdivided into different methods.

Manager class

Concurrency

Time progresses in the single `Clock` thread. All other threads wait for the `Clock` to reach a certain time, or for a certain number of minutes to elapse. This is accomplished through a synchronized list of `CountDownLatches`, counting down each minute until the waiting thread is released. The clock is also used to access a latch representing the start of the day (each thread waits for the main thread to release the latch once all threads are ready) and a barrier representing the end of the day (once each thread awaits the barrier, the day is over and stats can be calculated and displayed).

Standup meetings are handled using the reverse approach: While threads are blocked before the start of the day, unblocked during the day, and then blocked at the end, for meetings they run before the meeting, are blocked during the meeting, and are unblocked when it's over. Therefore a meeting starts with a `CyclicBarrier` with each participant reporting in before the meeting can start, and the meeting ends with a `CountDownLatch` with the participants blocked until the meeting is over.

Each team lead and manager has a lock ("busy") that must be acquired ("`getAttention()`") when asking a question. They must also acquire their own lock before going to a meeting or to lunch, to ensure they finish answering any question first. There is a second lock ("`canAnswerQuestions`") that team leads and the manager use to gain priority over the lock used for questions so they can go to meetings or lunch before being asked additional questions.

Alternatives

We experimented with using a `CyclicBarrier` within the `Clock` class to increment the minute only after each thread reports in (thus `Clock` wasn't a thread of its own, but a shared resource). This worked at first to ensure the threads are synced up throughout the day, but it precluded blocking the other threads while they're in meetings or asking questions, as they wouldn't be able to await the `CyclicBarrier` in the moments inbetween. `Phaser` could have resolved this, but we decided not to use it in order to maintain Java 6 compatibility. Instead, we reverted back to the original plan of using a thread.